Multiagent Reinforcement Learning: Rollout and Policy Iteration (Bertsekas, 2020)

Mikalai Korbit

IMT School for Advanced Studies Lucca

January 20, 2021

Outline

Sidenote – Path Integral Control

TODO

- TODO

Motivation

TODO

TODO.

Outline

TODO

TODO

Reinforcement Learning Problem

- Learning how to map situations to actions
- Trial-and-error search
- Delayed feedback
- Trade-off between exploration and exploitation
- Sequential decision making
- Agent's actions affect the subsequent data it receives

Reinforcement Learning Problem

- Learning how to map situations to actions
- Trial-and-error search
- Delayed feedback
- Trade-off between exploration and exploitation
- Sequential decision making
- Agent's actions affect the subsequent data it receives

Reinforcement Learning Problem

- Learning how to map situations to actions
- Trial-and-error search
- Delayed feedback
- Trade-off between exploration and exploitation
- Sequential decision making
- Agent's actions affect the subsequent data it receives

Conclusion

- RL methods can be applicable to a wide variety of problems
- Out-of-the-box models work but require fine-tuning and take longer to converge
- Simple methods like state discretization are worth exploring when training speed and solution complexity are of the essence

References


Reinforcement learning: an introduction, 2nd Edition. Richard S. Sutton, Andrew G. Barto.


Reinforcement learning lectures by David Silver. UCL. `http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html`


Playing Atari with Deep Reinforcement Learning. Mnih et al. `https://arxiv.org/abs/1312.5602`

Thanks for your attention!